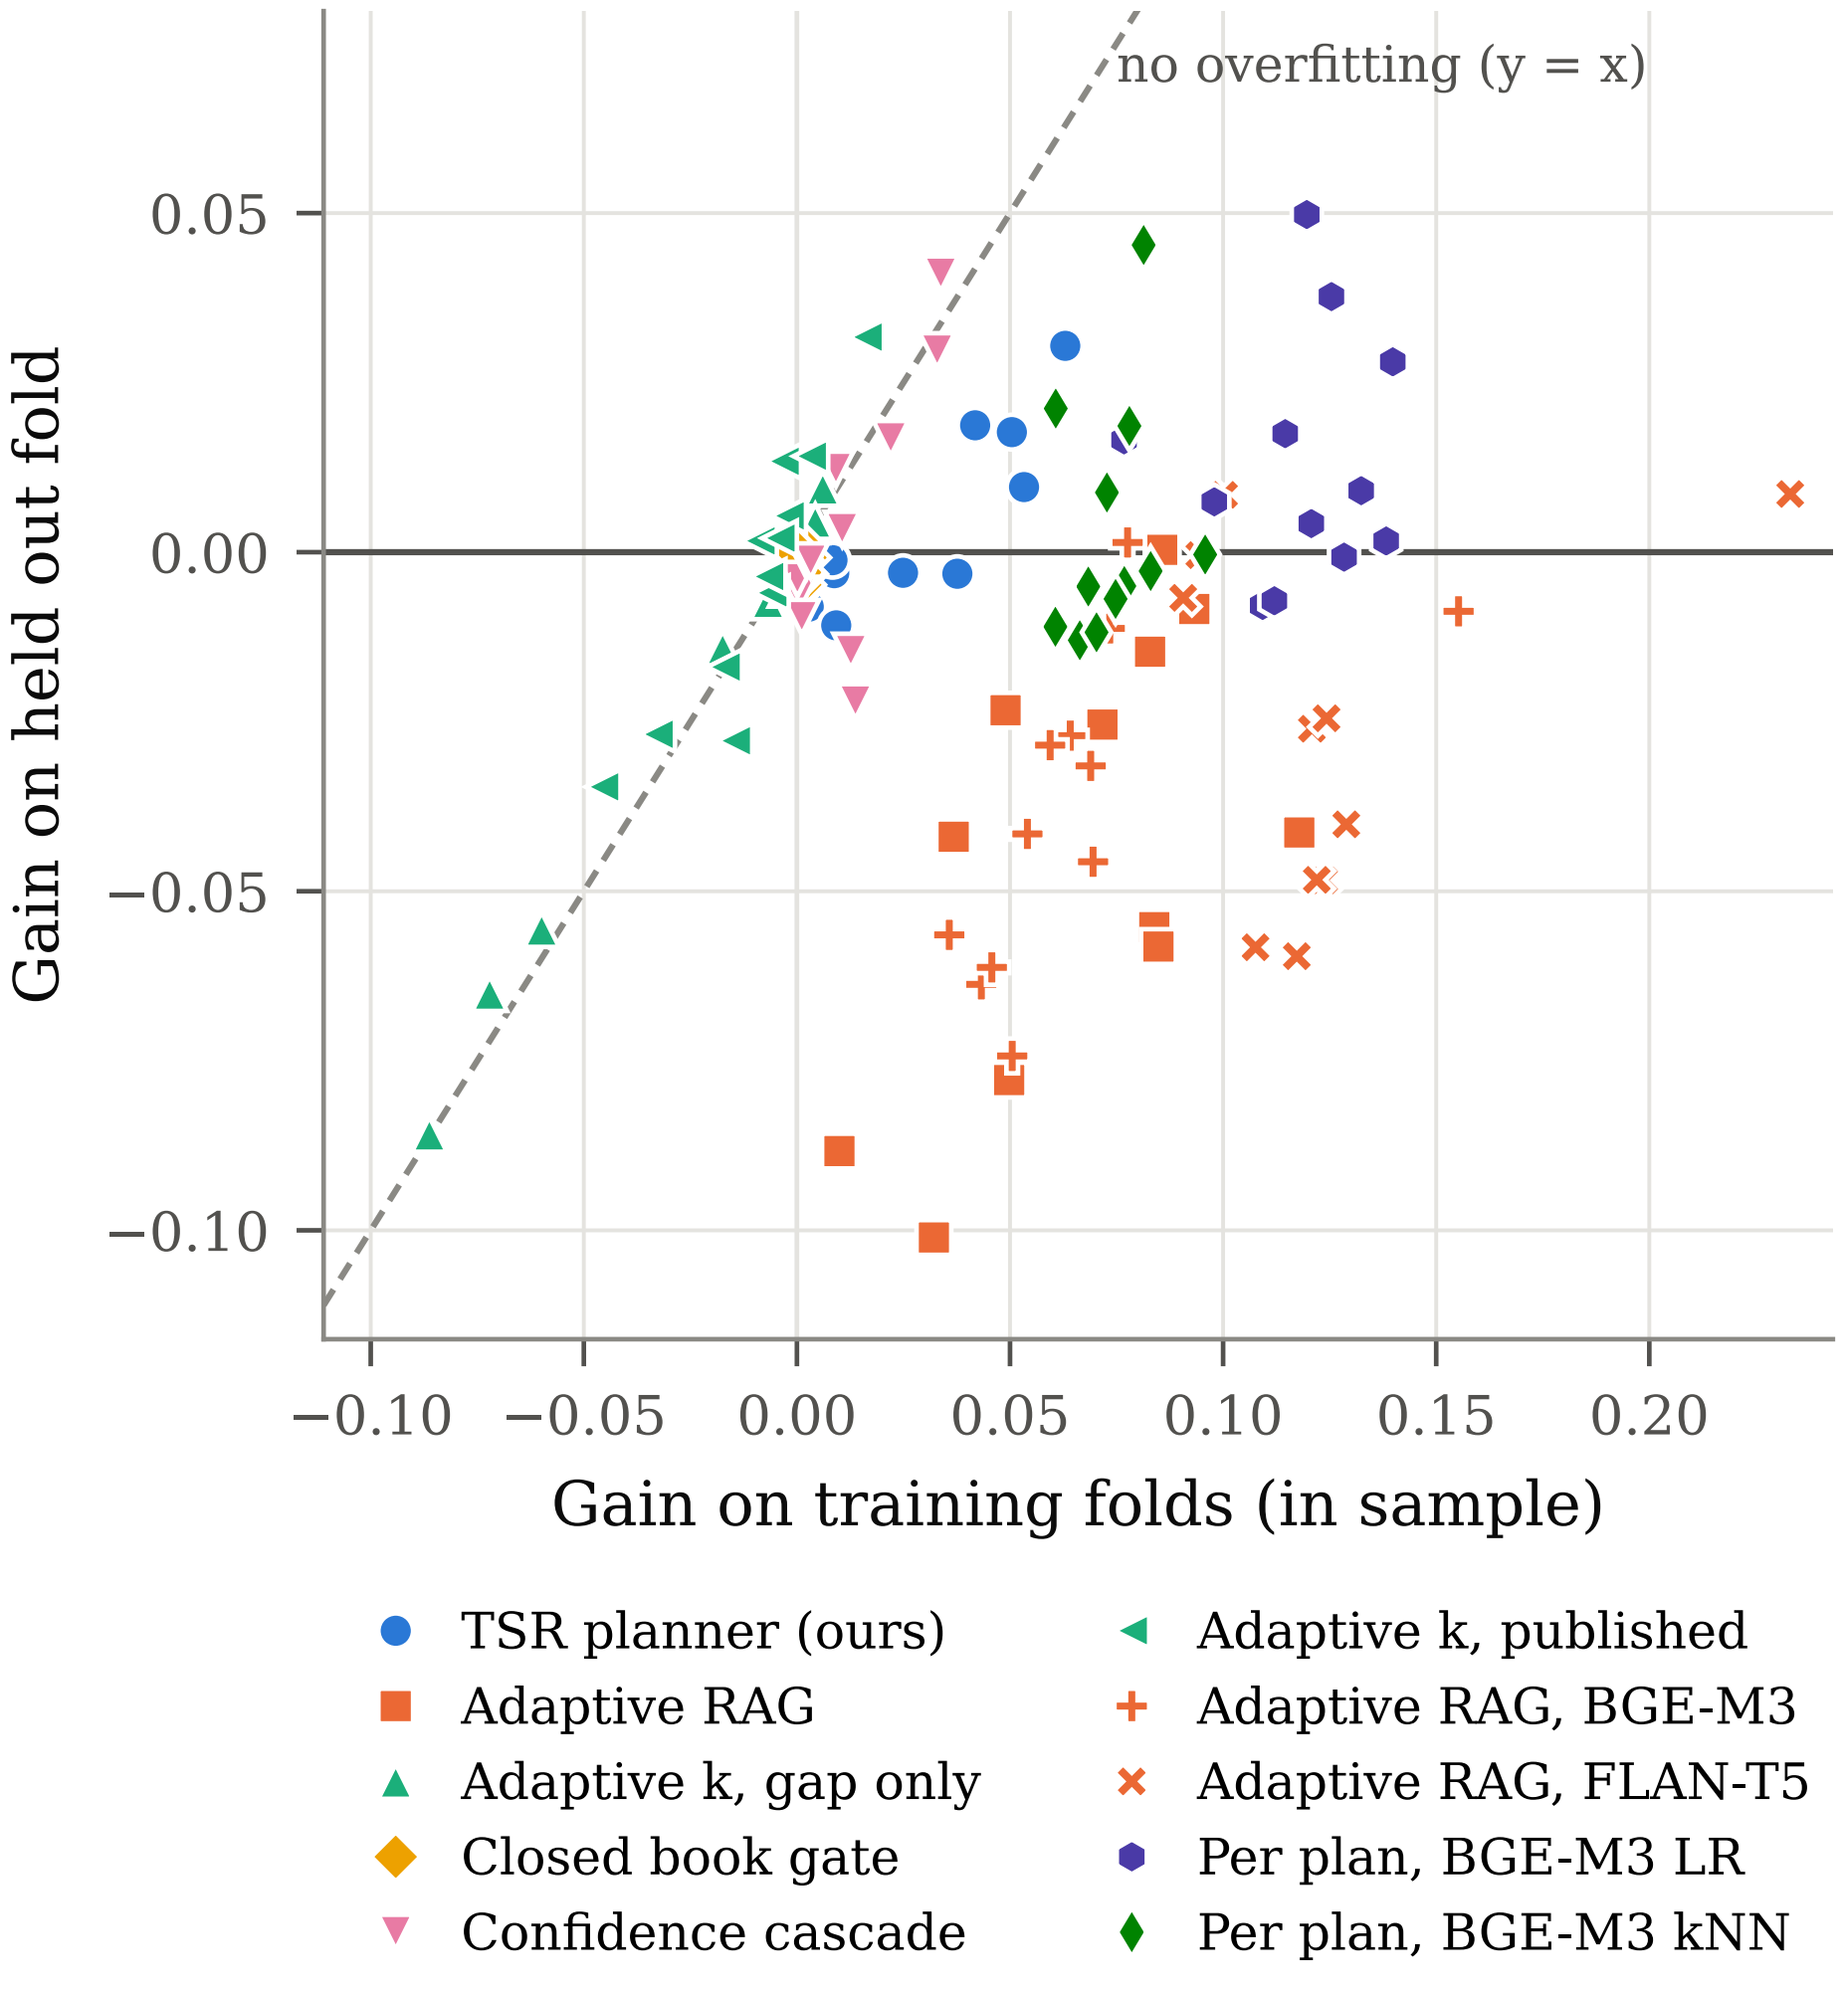
no overfitting (y = x)
0.05
0.00
−0.05
−0.10
Gain on held out fold
−0.10
−0.05
0.00
0.05
0.10
0.15
0.20
Gain on training folds (in sample)
TSR planner (ours)
Adaptive RAG
Adaptive k, gap only
Closed book gate
Confidence cascade
Adaptive k, published
Adaptive RAG, BGE-M3
Adaptive RAG, FLAN-T5
Per plan, BGE-M3 LR
Per plan, BGE-M3 kNN