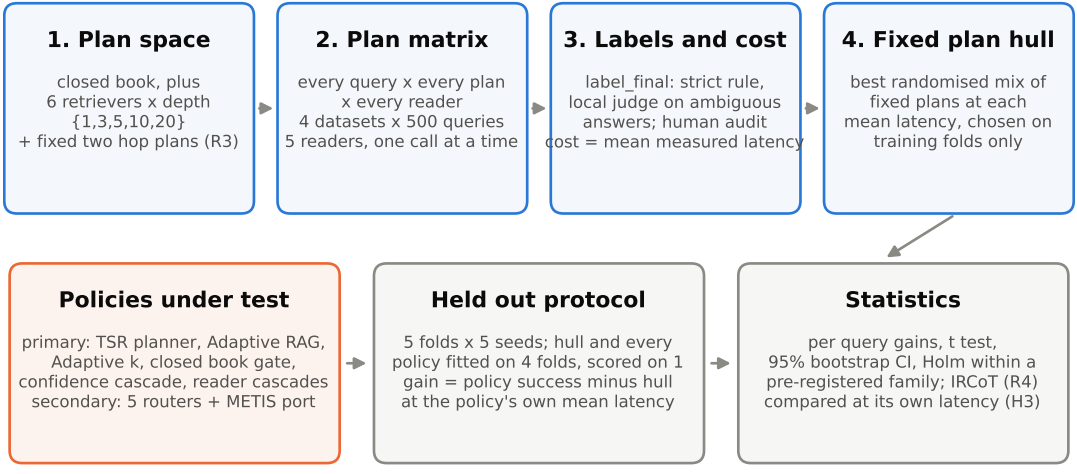
1. Plan space
closed book, plus
6 retrievers x depth
{1,3,5,10,20}
+ fixed two hop plans (R3)
2. Plan matrix
every query x every plan
x every reader
4 datasets x 500 queries
5 readers, one call at a time
3. Labels and cost
label_final: strict rule,
local judge on ambiguous
answers; human audit
cost = mean measured latency
4. Fixed plan hull
best randomised mix of
fixed plans at each
mean latency, chosen on
training folds only
Policies under test
primary: TSR planner, Adaptive RAG,
Adaptive k, closed book gate,
confidence cascade, reader cascades
secondary: 5 routers + METIS port
Held out protocol
5 folds x 5 seeds; hull and every
policy fitted on 4 folds, scored on 1
gain = policy success minus hull
at the policy's own mean latency
Statistics
per query gains, t test,
95% bootstrap CI, Holm within a
pre-registered family; IRCoT (R4)
compared at its own latency (H3)