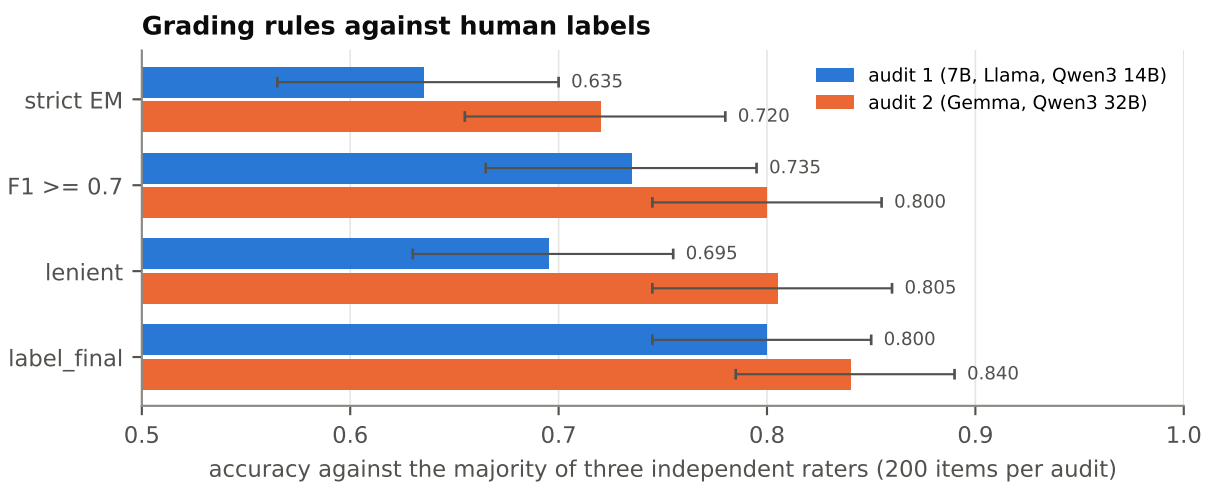
Grading rules against human labels
audit 1 (7B, Llama, Qwen3 14B)
audit 2 (Gemma, Qwen3 32B)
strict EM
F1 >= 0.7
lenient
label_final
0.635
0.720
0.735
0.800
0.695
0.805
0.800
0.840
0.5
0.6
0.7
0.8
0.9
1.0
accuracy against the majority of three independent raters (200 items per audit)